

Vulnerability Assessment Report: pygoat

Repository: [Akaolisangwu-projects/pygoat](https://github.com/Akaolisangwu-projects/pygoat)

Analyst: Akaolisa Ngwu

Summary

PyGoat is an intentionally vulnerable Django web application used for OWASP Top-10 training. Its purpose means many dangerous patterns exist by design and the repository must never be deployed as-is to production.

Key confirmed issues include:

- Public demo credentials published in the README.
- Docker/requirements files that allow trivial reproduction of the vulnerable environment.
- Use of dependency versions that are likely outdated.

These upgrades reduce container attack surface and demonstrate good security hygiene by integrating automated Snyk scans. I have already taken important remediation steps by merging two Snyk pull requests that upgrade the Docker base images to a more secure Python release, reducing OS- and runtime-level CVEs.

Changes Already Addressed (via Snyk)

Changes	From	To	Benefit
1	python:3.12.0a5-slim	python:3.14.0rc2-slim-trixie	Moves from an alpha image to a later patched release candidate, eliminating multiple Debian/CPython CVEs.
2	python:3.7.5-buster	python:3.14.0rc2-slim-trixie	Replaces an end-of-life 3.7 base (no fixes since 2023) with a maintained 3.14 release, removing known Debian-Buster vulnerabilities.

Confirmed Vulnerabilities & Risks

Category	Evidence	Risk	Recommended Remediation
Public demo credentials	README shows username: user / password: user12345	Critical	Remove or clearly sandbox these credentials; rotate any reused secrets immediately.
Reproducible vulnerable environment	Presence of requirements.txt, Dockerfile, and docker-compose.yml	High	Keep these files, but pin safe package versions and mark the repo as training only.
outdated Python dependencies	Requirements file exists but versions not audited	High	Run pip-audit/safety and update packages to patched versions.

Probable High-Risk Issues

These are common to intentionally vulnerable Django apps; confirm with static/dynamic scans:

- ❖ SQL Injection: raw SQL queries or unsafe `raw()` calls.
- ❖ Cross-Site Scripting (XSS): unescaped template output or use of `mark_safe`.
- ❖ CSRF Disabled: missing middleware or `{% csrf_token %}`.
- ❖ Weak Authentication: hard-coded or weak passwords, missing rate limits.
- ❖ Sensitive Data Exposure: debug mode enabled, hard-coded `SECRET_KEY`.
- ❖ Insecure Deserialization: use of `pickle` or unsafe YAML.
- ❖ Insufficient Logging & Monitoring.

Recommended Remediation Plan

Immediate (0-2 days)

1. Remove/rotate demo credentials in the README or ensure they point only to a hardened sandbox.
2. Clearly mark the repo and Docker images as *training only* to prevent accidental production deployment.
3. Set `DEBUG = False`, move `SECRET_KEY` to environment variables, and configure `ALLOWED_HOSTS`.

Short Term (1 week)

1. Run dependency and container scans
2. Patch or upgrade all high/critical CVEs found.
3. Add pre-commit hooks for secret detection (detect-secrets, git-secrets).

Medium Term (2–4 weeks)

1. Perform **static code analysis** (bandit -r .) and **dynamic scanning** (OWASP ZAP) to confirm SQLi/XSS/CSRF.
2. Harden Docker images further: use a non-root user, minimal base image, and multistage builds.
3. Implement CI/CD gates that fail on high severity findings.

Long Term

1. Maintain a secure production branch and keep intentionally vulnerable code only in a clearly labeled training branch.
2. Add structured logging and monitoring to detect exploitation attempts.

Key Hardening Checklist

- ☐ Remove credentials from README or rotate them.
- ☐ DEBUG=False in production settings.
- ☐ SECRET_KEY loaded from environment variable.
- ☐ SESSION_COOKIE_SECURE, CSRF_COOKIE_SECURE, SECURE_HSTS_SECONDS configured.
- ☐ Run pip-audit/safety regularly.
- ☐ Integrate Snyk scans into CI for both Python dependencies and Docker images.